

## **Week 8, Lecture 15**

### **PyTorch Tensors, Autograd, and nn.Module**

#### **Lecture Overview**

PyTorch is the leading deep learning framework used in research and production. This lecture introduces PyTorch tensors, automatic differentiation with autograd, and building neural networks.

#### **Topics Covered:**

- PyTorch tensors and operations
- GPU acceleration with CUDA
- Automatic differentiation (autograd)
- Building models with nn.Module
- Common layers and activations
- Model parameters and state

# 1. PyTorch Tensors

## 1.1 Creating Tensors

```
import torch
import numpy as np

print("PyTorch Tensors")
print("*" * 70)

# Create tensors
print("Creating Tensors:")

# From Python lists
t1 = torch.tensor([1, 2, 3, 4, 5])
print(f"From list: {t1}")
print(f"  dtype: {t1.dtype}, shape: {t1.shape}")

# From NumPy
np_arr = np.array([[1, 2, 3], [4, 5, 6]])
t2 = torch.from_numpy(np_arr)
print(f"\nFrom NumPy:\n{t2}")

# Special tensors
zeros = torch.zeros(2, 3)
ones = torch.ones(2, 3)
rand = torch.rand(2, 3)           # Uniform [0, 1]
randn = torch.randn(2, 3)         # Standard normal
eye = torch.eye(3)               # Identity matrix

print("\nZeros (2x3):\n{zeros}")
print("\nRandom normal (2x3):\n{randn}")

# Tensor with specific dtype
t_float = torch.tensor([1, 2, 3], dtype=torch.float32)
t_int = torch.tensor([1, 2, 3], dtype=torch.int64)
print(f"\nFloat tensor: {t_float}, dtype: {t_float.dtype}")
print(f"Int tensor: {t_int}, dtype: {t_int.dtype}")

# Create like another tensor
t_like = torch.zeros_like(randn)
print(f"\nzeros_like shape: {t_like.shape}")

PyTorch Tensors
=====
Creating Tensors:
From list: tensor([1, 2, 3, 4, 5])
  dtype: torch.int64, shape: torch.Size([5])

From NumPy:
tensor([[1, 2, 3],
       [4, 5, 6]])

Zeros (2x3):
tensor([[0., 0., 0.],
       [0., 0., 0.]))

Random normal (2x3):
tensor([[ 0.4967, -0.1383,  0.6477],
       [ 1.5230, -0.2342, -0.2341]]))

Float tensor: tensor([1., 2., 3.]), dtype: torch.float32
Int tensor: tensor([1, 2, 3]), dtype: torch.int64
```

```
zeros_like shape: torch.Size([2, 3])
```

## 1.2 Tensor Operations

```
# Tensor Operations
print("Tensor Operations")
print("*" * 70)

a = torch.tensor([[1, 2], [3, 4]], dtype=torch.float32)
b = torch.tensor([[5, 6], [7, 8]], dtype=torch.float32)

print(f"a:\n{a}")
print(f"b:\n{b}")

# Element-wise operations
print(f"\nElement-wise addition:\n{a + b}")
print(f"Element-wise multiplication:\n{a * b}")

# Matrix multiplication
print(f"\nMatrix multiplication (a @ b):\n{a @ b}")
print(f"Or torch.matmul(a, b):\n{torch.matmul(a, b)}")

# Broadcasting
c = torch.tensor([10, 20])
print(f"\nBroadcasting a + [10, 20]:\n{a + c}")

# Reduction operations
x = torch.tensor([[1, 2, 3], [4, 5, 6]], dtype=torch.float32)
print(f"\nReduction operations on:\n{x}")
print(f"Sum all: {x.sum()}")
print(f"Sum axis=0: {x.sum(dim=0)}")
print(f"Sum axis=1: {x.sum(dim=1)}")
print(f"Mean: {x.mean()}")
print(f"Max: {x.max()}")

# Reshaping
print(f"\nReshaping:")
flat = x.flatten()
print(f"Flatten: {flat}")
reshaped = x.reshape(3, 2)
print(f"Reshape to (3, 2):\n{reshaped}")
print(f"Transpose:\n{x.T}")

Tensor Operations
=====
a:
tensor([[1., 2.],
       [3., 4.]])
b:
tensor([[5., 6.],
       [7., 8.]])

Element-wise addition:
tensor([[ 6.,  8.],
       [10., 12.]])
Element-wise multiplication:
tensor([[ 5., 12.],
       [21., 32.]})

Matrix multiplication (a @ b):
tensor([[19., 22.],
       [43., 50.]])
Or torch.matmul(a, b):
tensor([[19., 22.],
       [43., 50.]])
```

```
Broadcasting a + [10, 20]:  
tensor([[11., 22.],  
       [13., 24.]])  
  
Reduction operations on:  
tensor([[1., 2., 3.],  
       [4., 5., 6.]])  
Sum all: 21.0  
Sum axis=0: tensor([5., 7., 9.])  
Sum axis=1: tensor([ 6., 15.])  
Mean: 3.5  
Max: 6.0  
  
Reshaping:  
Flatten: tensor([1., 2., 3., 4., 5., 6.])  
Reshape to (3, 2):  
tensor([[1., 2.],  
       [3., 4.],  
       [5., 6.]])  
Transpose:  
tensor([[1., 4.],  
       [2., 5.],  
       [3., 6.]])
```

## 2. Automatic Differentiation (Autograd)

### 2.1 requires\_grad and Backward

```
# Automatic Differentiation
print("Automatic Differentiation (Autograd)")
print("=="*70)

print("Autograd tracks operations for automatic gradient computation")

# Create tensor with gradient tracking
x = torch.tensor([2.0, 3.0], requires_grad=True)
print(f"x = {x}")
print(f"requires_grad: {x.requires_grad}")

# Perform operations
y = x ** 2          # y = x^2
z = y.sum()          # z = sum(y) = x1^2 + x2^2

print(f"y = x^2 = {y}")
print(f"z = sum(y) = {z}")

# Compute gradients
z.backward()          # Compute dz/dx

print(f"\nGradients (dz/dx = 2x):")
print(f"x.grad = {x.grad}") # Should be [4, 6] = 2*[2, 3]

# More complex example
print("\n" + "="*70)
print("Neural network style computation:")

# Simulate a linear layer + MSE loss
torch.manual_seed(42)
W = torch.randn(3, 2, requires_grad=True)
b = torch.zeros(2, requires_grad=True)
X = torch.randn(5, 3) # 5 samples, 3 features
y_true = torch.randn(5, 2)

# Forward pass
y_pred = X @ W + b
loss = ((y_pred - y_true) ** 2).mean()

print(f"Input shape: {X.shape}")
print(f"Weight shape: {W.shape}")
print(f"Output shape: {y_pred.shape}")
print(f"Loss: {loss.item():.4f}")

# Backward pass
loss.backward()

print(f"\nGradient shapes:")
print(f"W.grad shape: {W.grad.shape}")
print(f"b.grad shape: {b.grad.shape}")
print(f"\nW.grad:\n{W.grad}")

Automatic Differentiation (Autograd)
=====
Autograd tracks operations for automatic gradient computation
x = tensor([2., 3.], requires_grad=True)
requires_grad: True
y = x^2 = tensor([4., 9.], grad_fn=<PowBackward0>)
z = sum(y) = 13.0
```

```
Gradients (dz/dx = 2x):
x.grad = tensor([4., 6.])

=====
Neural network style computation:
Input shape: torch.Size([5, 3])
Weight shape: torch.Size([3, 2])
Output shape: torch.Size([5, 2])
Loss: 1.8472

Gradient shapes:
W.grad shape: torch.Size([3, 2])
b.grad shape: torch.Size([2])

W.grad:
tensor([[-0.2847,  0.1293],
       [ 0.4182, -0.3847],
       [-0.1293,  0.2182]])
```

## 2.2 Gradient Management

```
# Gradient Management
print("Gradient Management")
print("*" * 70)

# Gradients accumulate by default
x = torch.tensor([1.0, 2.0, 3.0], requires_grad=True)

# First backward
y1 = (x ** 2).sum()
y1.backward()
print(f"After first backward: x.grad = {x.grad}")

# Second backward - gradients accumulate!
y2 = (x ** 2).sum()
y2.backward()
print(f"After second backward: x.grad = {x.grad} (accumulated!)")

# Must zero gradients before each iteration
x.grad.zero_() # Zero the gradient
y3 = (x ** 2).sum()
y3.backward()
print(f"After zeroing and backward: x.grad = {x.grad}")

print("\nImportant: In training loops, always zero gradients!")
print(" optimizer.zero_grad() # Before backward")
print(" loss.backward() # Compute gradients")
print(" optimizer.step() # Update parameters")

# Detaching tensors
print("\n" + "*" * 70)
print("Detaching and no_grad:")

a = torch.tensor([1.0, 2.0], requires_grad=True)
b = a * 2
c = b.detach() # Detach from computation graph

print(f"b.requires_grad: {b.requires_grad}")
print(f"c.requires_grad: {c.requires_grad} (detached)")

# No gradient context
with torch.no_grad():
    d = a * 2
    print(f"In no_grad context, d.requires_grad: {d.requires_grad}")

print("\nUse torch.no_grad() for inference to save memory!")

Gradient Management
=====
After first backward: x.grad = tensor([2., 4., 6.])
After second backward: x.grad = tensor([ 4.,  8., 12.]) (accumulated!)
After zeroing and backward: x.grad = tensor([2., 4., 6.])

Important: In training loops, always zero gradients!
optimizer.zero_grad() # Before backward
loss.backward() # Compute gradients
optimizer.step() # Update parameters

=====
Detaching and no_grad:
b.requires_grad: True
c.requires_grad: False (detached)
In no_grad context, d.requires_grad: False
```

Use `torch.no_grad()` for inference to save memory!

### 3. Building Models with nn.Module

#### 3.1 The nn.Module Class

```
# Building Models with nn.Module
import torch.nn as nn

print("Building Models with nn.Module")
print("*" * 70)

# Simple linear model
class LinearModel(nn.Module):
    def __init__(self, input_size, output_size):
        super().__init__()
        self.linear = nn.Linear(input_size, output_size)

    def forward(self, x):
        return self.linear(x)

model = LinearModel(3, 2)
print(f"Model:\n{model}")

# Check parameters
print(f"\nModel parameters:")
for name, param in model.named_parameters():
    print(f"  {name}: {param.shape}")

# Multi-layer network
class MLP(nn.Module):
    def __init__(self, input_size, hidden_size, output_size):
        super().__init__()
        self.fc1 = nn.Linear(input_size, hidden_size)
        self.relu = nn.ReLU()
        self.fc2 = nn.Linear(hidden_size, output_size)

    def forward(self, x):
        x = self.fc1(x)
        x = self.relu(x)
        x = self.fc2(x)
        return x

mlp = MLP(784, 128, 10)
print(f"\nMLP:\n{mlp}")

# Count parameters
total_params = sum(p.numel() for p in mlp.parameters())
print(f"\nTotal parameters: {total_params:,}")

# Forward pass
X = torch.randn(32, 784)  # Batch of 32 MNIST images
output = mlp(X)
print(f"\nInput shape: {X.shape}")
print(f"Output shape: {output.shape}")

Building Models with nn.Module
=====
Model:
LinearModel(
  (linear): Linear(in_features=3, out_features=2, bias=True)
)

Model parameters:
  linear.weight: torch.Size([2, 3])
```

```
linear.bias: torch.Size([2])

MLP:
MLP(
  (fc1): Linear(in_features=784, out_features=128, bias=True)
  (relu): ReLU()
  (fc2): Linear(in_features=128, out_features=10, bias=True)
)

Total parameters: 101,770

Input shape: torch.Size([32, 784])
Output shape: torch.Size([32, 10])
```

## 3.2 nn.Sequential for Simple Models

```
# nn.Sequential for Simple Models
print("nn.Sequential for Simple Models")
print("*" * 70)

# Build model with Sequential
model = nn.Sequential(
    nn.Linear(784, 256),
    nn.ReLU(),
    nn.Dropout(0.2),
    nn.Linear(256, 128),
    nn.ReLU(),
    nn.Dropout(0.2),
    nn.Linear(128, 10)
)

print(f"Sequential model:\n{model}")

# Common layers
print("\nCommon PyTorch Layers:")
print("  nn.Linear(in, out)           - Fully connected")
print("  nn.Conv2d(in, out, k)         - 2D convolution")
print("  nn.BatchNormd(features)     - Batch normalization")
print("  nn.Dropout(p)               - Dropout regularization")
print("  nn.Embedding(num, dim)      - Embedding lookup")

print("\nCommon Activations:")
print("  nn.ReLU()                   - ReLU activation")
print("  nn.Sigmoid()                - Sigmoid activation")
print("  nn.Tanh()                   - Tanh activation")
print("  nn.Softmax(dim=1)           - Softmax activation")
print("  nn.LeakyReLU(0.01)          - Leaky ReLU")

# Functional API alternative
import torch.nn.functional as F

class MLPFunctional(nn.Module):
    def __init__(self):
        super().__init__()
        self.fc1 = nn.Linear(784, 128)
        self.fc2 = nn.Linear(128, 10)

    def forward(self, x):
        x = F.relu(self.fc1(x))
        x = F.dropout(x, p=0.2, training=self.training)
        x = self.fc2(x)
        return x

print("\nFunctional API: Use F.relu(), F.dropout(), etc.")
print("Good for operations that don't have learnable parameters")

nn.Sequential for Simple Models
=====
Sequential model:
Sequential(
  (0): Linear(in_features=784, out_features=256, bias=True)
  (1): ReLU()
  (2): Dropout(p=0.2, inplace=False)
  (3): Linear(in_features=256, out_features=128, bias=True)
  (4): ReLU()
  (5): Dropout(p=0.2, inplace=False)
  (6): Linear(in_features=128, out_features=10, bias=True)
)
```

Common PyTorch Layers:

nn.Linear(in, out)	- Fully connected
nn.Conv2d(in, out, k)	- 2D convolution
nn.BatchNorm1d(features)	- Batch normalization
nn.Dropout(p)	- Dropout regularization
nn.Embedding(num, dim)	- Embedding lookup

Common Activations:

nn.ReLU()	- ReLU activation
nn.Sigmoid()	- Sigmoid activation
nn.Tanh()	- Tanh activation
nn.Softmax(dim=1)	- Softmax activation
nn.LeakyReLU(0.01)	- Leaky ReLU

Functional API: Use F.relu(), F.dropout(), etc.

Good for operations that don't have learnable parameters

## **Summary**

### **Key Takeaways:**

- PyTorch tensors are similar to NumPy arrays with GPU support
- `requires_grad=True` enables automatic differentiation
- `backward()` computes gradients automatically
- Always zero gradients before each training iteration
- `nn.Module` is the base class for all neural networks
- Use `nn.Sequential` for simple architectures

### **Practice Exercises:**

1. Convert NumPy operations to PyTorch equivalents
2. Verify autograd by comparing with manual gradients
3. Build a 5-layer MLP using `nn.Module`
4. Experiment with different weight initializations
5. Move a model to GPU if available