

Pooling, Feature Hierarchies, and Classic Architectures

Lecture Overview

This lecture covers pooling operations, CNN feature hierarchies, and classic architectures that shaped deep learning.

Topics Covered:

- Max and average pooling
- Global average pooling
- Feature hierarchy in CNNs
- LeNet-5, AlexNet, VGGNet

1. Pooling Layers

1.1 Max and Average Pooling

```
import torch
import torch.nn as nn

print("Pooling Layers")
print("*" * 70)

print("Pooling: Reduce spatial dimensions")
print(" - Provides translation invariance")
print(" - Reduces computation")
print(" - Increases receptive field")

x = torch.tensor([[[[1, 2, 3, 4],
                   [5, 6, 7, 8],
                   [9, 10, 11, 12],
                   [13, 14, 15, 16]]]], dtype=torch.float32)

print(f"\nInput (4x4):\n{x.squeeze()}\n")

max_pool = nn.MaxPool2d(kernel_size=2, stride=2)
avg_pool = nn.AvgPool2d(kernel_size=2, stride=2)

print(f"\nMaxPool2d(2x2):\n{max_pool(x).squeeze()}\n")
print(" Takes maximum in each region")

print(f"\nAvgPool2d(2x2):\n{avg_pool(x).squeeze()}\n")
print(" Takes average in each region")

print("\nMax vs Average:")
print(" MaxPool: Preserves strongest activations")
print(" AvgPool: Smoother, sometimes at end of network")

Pooling Layers
=====
Pooling: Reduce spatial dimensions
- Provides translation invariance
- Reduces computation
- Increases receptive field

Input (4x4):
tensor([[ 1.,  2.,  3.,  4.],
       [ 5.,  6.,  7.,  8.],
       [ 9., 10., 11., 12.],
       [13., 14., 15., 16.]])\n\n

MaxPool2d(2x2):
tensor([[ 6.,  8.],
       [14., 16.]])
Takes maximum in each region\n\n

AvgPool2d(2x2):
tensor([[ 3.5,  5.5],
       [11.5, 13.5]])
Takes average in each region\n\n

Max vs Average:
MaxPool: Preserves strongest activations
AvgPool: Smoother, sometimes at end of network
```

1.2 Global Average Pooling

```
print("Global Average Pooling (GAP)")
print("*" * 70)

print("GAP: Reduces each feature map to single value")
print(" Alternative to flatten + FC layers")
print(" Much fewer parameters!")

x = torch.randn(4, 512, 7, 7)
print(f"\nInput: {x.shape}")
```

```
gap = nn.AdaptiveAvgPool2d(1)
out = gap(x).view(x.size(0), -1)
print(f"After GAP: {out.shape}")

print("\nParameter Comparison:")
fc_params = 512 * 7 * 7 * 1000
gap_params = 512 * 1000
print(f"  Flatten + FC: {fc_params:,}")
print(f"  GAP + FC: {gap_params:,}")
print(f"  Reduction: {fc_params/gap_params:.0f}x")

Global Average Pooling (GAP)
=====
GAP: Reduces each feature map to single value
Alternative to flatten + FC layers
Much fewer parameters!

Input: torch.Size([4, 512, 7, 7])
After GAP: torch.Size([4, 512])

Parameter Comparison:
  Flatten + FC: 25,088,000
  GAP + FC: 512,000
  Reduction: 49x
```

2. Feature Hierarchies

2.1 What CNNs Learn

```
print("Feature Hierarchies in CNNs")
print("="*70)

print("CNNs learn hierarchical representations:")

print("\nLayer 1 (Early):")
print(" - Edges at various orientations")
print(" - Color blobs")
print(" - Simple textures")

print("\nLayers 2-3 (Middle):")
print(" - Corners and contours")
print(" - Texture patterns")
print(" - Simple shapes")

print("\nLayers 4-5 (Late):")
print(" - Object parts (eyes, wheels)")
print(" - Complex patterns")

print("\nFinal Layers:")
print(" - Whole objects")
print(" - Scene categories")

class VGGLike(nn.Module):
    def __init__(self):
        super().__init__()
        self.block1 = nn.Sequential(
            nn.Conv2d(3, 64, 3, padding=1), nn.ReLU(), nn.MaxPool2d(2))
        self.block2 = nn.Sequential(
            nn.Conv2d(64, 128, 3, padding=1), nn.ReLU(), nn.MaxPool2d(2))
        self.block3 = nn.Sequential(
            nn.Conv2d(128, 256, 3, padding=1), nn.ReLU(), nn.MaxPool2d(2))

    def forward(self, x):
        f1 = self.block1(x)
        f2 = self.block2(f1)
        f3 = self.block3(f2)
        return f1, f2, f3

model = VGGLike()
x = torch.randn(1, 3, 64, 64)
f1, f2, f3 = model(x)

print(f"\nFeature map shapes:")
print(f" Input: {x.shape}")
print(f" Block 1: {f1.shape} (edges)")
print(f" Block 2: {f2.shape} (textures)")
print(f" Block 3: {f3.shape} (parts)")

Feature Hierarchies in CNNs
=====
CNNs learn hierarchical representations:

Layer 1 (Early):
- Edges at various orientations
- Color blobs
- Simple textures

Layers 2-3 (Middle):
- Corners and contours
- Texture patterns
- Simple shapes

Layers 4-5 (Late):
- Object parts (eyes, wheels)
- Complex patterns

Final Layers:
- Whole objects
- Scene categories
```

```
Feature map shapes:  
Input: torch.Size([1, 3, 64, 64])  
Block 1: torch.Size([1, 64, 32, 32]) (edges)  
Block 2: torch.Size([1, 128, 16, 16]) (textures)  
Block 3: torch.Size([1, 256, 8, 8]) (parts)
```

3. Classic CNN Architectures

3.1 LeNet-5

```
print("LeNet-5 (LeCun et al., 1998)")
print("*" * 70)

print("First successful CNN - digit recognition")

class LeNet5(nn.Module):
    def __init__(self, num_classes=10):
        super().__init__()
        self.features = nn.Sequential(
            nn.Conv2d(1, 6, 5), nn.ReLU(), nn.AvgPool2d(2),
            nn.Conv2d(6, 16, 5), nn.ReLU(), nn.AvgPool2d(2),
        )
        self.classifier = nn.Sequential(
            nn.Flatten(),
            nn.Linear(16*5*5, 120), nn.ReLU(),
            nn.Linear(120, 84), nn.ReLU(),
            nn.Linear(84, num_classes)
        )
    def forward(self, x):
        return self.classifier(self.features(x))

model = LeNet5()
params = sum(p.numel() for p in model.parameters())
print(f"Parameters: {params:,}")

print("\nKey Contributions:")
print("- Local receptive fields")
print("- Weight sharing")
print("- Subsampling (pooling)")

LeNet-5 (LeCun et al., 1998)
=====
First successful CNN - digit recognition
Parameters: 61,706

Key Contributions:
- Local receptive fields
- Weight sharing
- Subsampling (pooling)
```

3.2 AlexNet

```
print("AlexNet (Krizhevsky et al., 2012)")
print("*" * 70)

print("Won ImageNet 2012 - started deep learning revolution")

print("\nArchitecture (simplified):")
print("Conv(3->96, 11x11, stride=4) -> ReLU -> MaxPool")
print("Conv(96->256, 5x5) -> ReLU -> MaxPool")
print("Conv(256->384, 3x3) -> ReLU")
print("Conv(384->384, 3x3) -> ReLU")
print("Conv(384->256, 3x3) -> ReLU -> MaxPool")
print("FC(9216->4096) -> FC(4096->4096) -> FC(4096->1000)")

print("\nParameters: ~61M")

print("\nKey Innovations:")
print("- First to use ReLU activation")
print("- Dropout for regularization")
print("- Heavy data augmentation")
print("- GPU training (2 GPUs)")
print("- Local Response Normalization")

AlexNet (Krizhevsky et al., 2012)
=====
Won ImageNet 2012 - started deep learning revolution

Architecture (simplified):
```

```

Conv(3->96, 11x11, stride=4) -> ReLU -> MaxPool
Conv(96->256, 5x5) -> ReLU -> MaxPool
Conv(256->384, 3x3) -> ReLU
Conv(384->384, 3x3) -> ReLU
Conv(384->256, 3x3) -> ReLU -> MaxPool
FC(9216->4096) -> FC(4096->4096) -> FC(4096->1000)

```

Parameters: ~61M

Key Innovations:

- First to use ReLU activation
- Dropout for regularization
- Heavy data augmentation
- GPU training (2 GPUs)
- Local Response Normalization

3.3 VGGNet

```

print("VGGNet (Simonyan & Zisserman, 2014)")
print("*" * 70)

print("Key insight: Use small 3x3 filters, stack deep")
print(" Two 3x3 = same receptive field as one 5x5")
print(" Three 3x3 = one 7x7")
print(" But fewer params and more nonlinearities!")

one_7x7 = 7 * 7 * 64 * 64
three_3x3 = 3 * (3 * 3 * 64 * 64)
print(f"\nParameter comparison:")
print(f" One 7x7: {one_7x7:,}")
print(f" Three 3x3: {three_3x3:,}")

print("\nVGG16 Architecture:")
print(" 64, 64, pool")
print(" 128, 128, pool")
print(" 256, 256, 256, pool")
print(" 512, 512, 512, pool")
print(" 512, 512, 512, pool")
print(" FC 4096, FC 4096, FC 1000")

print("\nVGG16 Parameters: ~138M")

print("\nArchitecture Comparison:")
print(f"{'Model':<10} {'Year':>6} {'Depth':>6} {'Params':>10}")
print("-" * 35)
print(f"{'LeNet':<10} {'1998':>6} {'5':>6} {'60K':>10}")
print(f"{'AlexNet':<10} {'2012':>6} {'8':>6} {'61M':>10}")
print(f"{'VGG16':<10} {'2014':>6} {'16':>6} {'138M':>10}")

VGGNet (Simonyan & Zisserman, 2014)
=====
Key insight: Use small 3x3 filters, stack deep
Two 3x3 = same receptive field as one 5x5
Three 3x3 = one 7x7
But fewer params and more nonlinearities!

Parameter comparison:
One 7x7: 200,704
Three 3x3: 110,592

VGG16 Architecture:
64, 64, pool
128, 128, pool
256, 256, 256, pool
512, 512, 512, pool
512, 512, 512, pool
FC 4096, FC 4096, FC 1000

VGG16 Parameters: ~138M

Architecture Comparison:
Model      Year   Depth     Params
-----
LeNet       1998      5      60K
AlexNet    2012      8      61M

```

Summary

Key Takeaways:

- Pooling reduces spatial dimensions and provides invariance
- GAP reduces parameters vs flatten+FC
- CNNs learn hierarchical features: edges -> textures -> parts -> objects
- LeNet proved CNNs work
- AlexNet started deep learning revolution
- VGG showed small 3x3 filters stacked deep are effective

Practice Exercises:

1. Implement max pooling from scratch
2. Implement LeNet-5 and train on MNIST
3. Calculate receptive fields
4. Visualize feature maps at different depths
5. Compare different pooling strategies