

## **Modern Architectures: ResNet and EfficientNet**

### **Lecture Overview**

This lecture covers modern CNN architectures that enable training very deep networks through residual connections and efficient scaling.

Topics Covered:

- The degradation problem
- ResNet skip connections
- BasicBlock and Bottleneck designs
- EfficientNet compound scaling
- MBConv and squeeze-excitation

# 1. The Degradation Problem

## 1.1 Why Deeper Is Not Always Better

```
import torch
import torch.nn as nn
import numpy as np

print("The Degradation Problem")
print("=*70)

print("Observation:")
print("  56-layer network has HIGHER training error than 20-layer!")
print("  Not overfitting - training error is worse")

print("\nKey insight (He et al., 2015):")
print("  If deeper layers were identity mappings,")
print("  deeper network should be &gt;= as good as shallow")
print("  But plain networks struggle to learn identity!")

print("\nSolution: Residual connections (skip connections)")
print("  Plain:    y = F(x)")
print("  Residual: y = F(x) + x"

print("\nThe network learns F(x) = H(x) - x (the residual)")
print("If identity is optimal, F(x) learns to be zero")
print("Learning zero is easier than learning identity!")

The Degradation Problem
=====
Observation:
  56-layer network has HIGHER training error than 20-layer!
  Not overfitting - training error is worse

Key insight (He et al., 2015):
  If deeper layers were identity mappings,
  deeper network should be &gt;= as good as shallow
  But plain networks struggle to learn identity!

Solution: Residual connections (skip connections)
  Plain:    y = F(x)
  Residual: y = F(x) + x

The network learns F(x) = H(x) - x (the residual)
If identity is optimal, F(x) learns to be zero
Learning zero is easier than learning identity!
```

## 2. ResNet Architecture

### 2.1 Basic Block

```
print("ResNet Basic Block")
print("*" * 70)

class BasicBlock(nn.Module):
    """For ResNet-18, 34"""
    def __init__(self, in_ch, out_ch, stride=1):
        super().__init__()
        self.conv1 = nn.Conv2d(in_ch, out_ch, 3, stride=stride, padding=1, bias=False)
        self.bn1 = nn.BatchNorm2d(out_ch)
        self.conv2 = nn.Conv2d(out_ch, out_ch, 3, padding=1, bias=False)
        self.bn2 = nn.BatchNorm2d(out_ch)

        self.shortcut = nn.Sequential()
        if stride != 1 or in_ch != out_ch:
            self.shortcut = nn.Sequential(
                nn.Conv2d(in_ch, out_ch, 1, stride=stride, bias=False),
                nn.BatchNorm2d(out_ch)
            )

    def forward(self, x):
        identity = x
        out = torch.relu(self.bn1(self.conv1(x)))
        out = self.bn2(self.conv2(out))
        out += self.shortcut(identity) # Skip connection!
        return torch.relu(out)

block = BasicBlock(64, 64)
x = torch.randn(1, 64, 32, 32)
print(f"Input: {x.shape}")
print(f"Output: {block(x).shape}")

print("\nGradient flow through residual:")
print("dy/dx = dF/dx + 1")
print("The '+1' ensures gradient always flows!")

ResNet Basic Block
=====
Input: torch.Size([1, 64, 32, 32])
Output: torch.Size([1, 64, 32, 32])

Gradient flow through residual:
dy/dx = dF/dx + 1
The '+1' ensures gradient always flows!
```

### 2.2 Bottleneck Block

```
print("Bottleneck Block (ResNet-50, 101, 152)")
print("*" * 70)

print("Design: 1x1 reduce -> 3x3 process -> 1x1 expand")

class Bottleneck(nn.Module):
    expansion = 4

    def __init__(self, in_ch, mid_ch, stride=1):
        super().__init__()
        out_ch = mid_ch * self.expansion
        self.conv1 = nn.Conv2d(in_ch, mid_ch, 1, bias=False)
        self.bn1 = nn.BatchNorm2d(mid_ch)
        self.conv2 = nn.Conv2d(mid_ch, mid_ch, 3, stride=stride, padding=1, bias=False)
        self.bn2 = nn.BatchNorm2d(mid_ch)
        self.conv3 = nn.Conv2d(mid_ch, out_ch, 1, bias=False)
        self.bn3 = nn.BatchNorm2d(out_ch)

        self.shortcut = nn.Sequential()
        if stride != 1 or in_ch != out_ch:
            self.shortcut = nn.Sequential(
                nn.Conv2d(in_ch, out_ch, 1, stride=stride, bias=False),
                nn.BatchNorm2d(out_ch)
```

```

    )

def forward(self, x):
    identity = x
    out = torch.relu(self.bn1(self.conv1(x)))
    out = torch.relu(self.bn2(self.conv2(out)))
    out = self.bn3(self.conv3(out))
    out += self.shortcut(identity)
    return torch.relu(out)

# Parameter comparison
basic = 2 * (3*3*256*256)
bottleneck = (1*1*256*64) + (3*3*64*64) + (1*1*64*256)
print(f"\nParams (256 channels):")
print(f"  BasicBlock: {basic:.1f}")
print(f"  Bottleneck: {bottleneck:.1f}")
print(f"  Reduction: {basic/bottleneck:.1f}x")

Bottleneck Block (ResNet-50, 101, 152)
=====
Design: 1x1 reduce -> 3x3 process -> 1x1 expand

Params (256 channels):
  BasicBlock: 1,179,648
  Bottleneck: 69,632
  Reduction: 16.9x

```

## 2.3 ResNet Variants

```

print("ResNet Variants")
print("*" * 70)

print(f"{'Model':<12} {'Blocks':>20} {'Params':>10}")
print("-" * 45)
print(f"{'ResNet-18':<12} {[2, 2, 2, 2]}:{>20} {'11.7M':>10}")
print(f"{'ResNet-34':<12} {[3, 4, 6, 3]}:{>20} {'21.8M':>10}")
print(f"{'ResNet-50':<12} {[3, 4, 6, 3]*}:{>20} {'25.6M':>10}")
print(f"{'ResNet-101':<12} {[3, 4, 23, 3]*}:{>20} {'44.5M':>10}")
print(f"{'ResNet-152':<12} {[3, 8, 36, 3]*}:{>20} {'60.2M':>10}")
print("* uses Bottleneck blocks")

print("\nUsing pre-trained ResNet:")
print("  from torchvision import models")
print("  model = models.resnet50(pretrained=True)")
print("  model.fc = nn.Linear(2048, num_classes)")

ResNet Variants
=====
Model      Blocks      Params
-----
ResNet-18      [2, 2, 2, 2]      11.7M
ResNet-34      [3, 4, 6, 3]      21.8M
ResNet-50      [3, 4, 6, 3]*     25.6M
ResNet-101     [3, 4, 23, 3]*    44.5M
ResNet-152     [3, 8, 36, 3]*    60.2M
* uses Bottleneck blocks

Using pre-trained ResNet:
from torchvision import models
model = models.resnet50(pretrained=True)
model.fc = nn.Linear(2048, num_classes)

```

## 3. EfficientNet

### 3.1 Compound Scaling

```
print("EfficientNet (Tan & Le, 2019)")  
print("*" * 70)  
  
print("Key insight: Scale depth, width, resolution together")  
  
print("\nScaling dimensions:")  
print("  Depth: Number of layers")  
print("  Width: Number of channels")  
print("  Resolution: Input image size")  
  
print("\nCompound scaling:")  
print("  depth = alpha^phi")  
print("  width = beta^phi")  
print("  resolution = gamma^phi")  
print("  Constraint: alpha * beta^2 * gamma^2 = 2")  
  
print("\nEfficientNet Family:")  
print(f"{'Model':<8} {'Params':>8} {'Top-1':>8}")  
print("-" * 26)  
for m, p, a in [("B0", "5.3M", "77.1%"), ("B1", "7.8M", "79.1%"),  
                ("B2", "9.2M", "80.1%"), ("B3", "12M", "81.6%"),  
                ("B4", "19M", "82.9%"), ("B7", "66M", "84.3%")]:  
    print(f"{m:<8} {p:>8} {a:>8}")  
  
print("\nComparison:")  
print("  ResNet-50: 25.6M params, 76% Top-1")  
print("  EfficientNet-B0: 5.3M params, 77% Top-1")  
print("  5x fewer params, better accuracy!")  
  
EfficientNet (Tan & Le, 2019)  
=====
```

Key insight: Scale depth, width, resolution together

Scaling dimensions:

Depth: Number of layers  
Width: Number of channels  
Resolution: Input image size

Compound scaling:

depth = alpha^phi  
width = beta^phi  
resolution = gamma^phi  
Constraint: alpha \* beta^2 \* gamma^2 = 2

EfficientNet Family:

Model	Params	Top-1
B0	5.3M	77.1%
B1	7.8M	79.1%
B2	9.2M	80.1%
B3	12M	81.6%
B4	19M	82.9%
B7	66M	84.3%

Comparison:

ResNet-50: 25.6M params, 76% Top-1  
EfficientNet-B0: 5.3M params, 77% Top-1  
5x fewer params, better accuracy!

### 3.2 MBConv Block

```
print("MBConv (Mobile Inverted Bottleneck)")  
print("*" * 70)  
  
print("MBConv structure:")  
print("  1. Expand: 1x1 conv (increase channels)")  
print("  2. Depthwise: 3x3 or 5x5 depthwise conv")  
print("  3. Squeeze-Excitation: Channel attention")  
print("  4. Project: 1x1 conv (reduce channels)")
```

```

class SqueezeExcitation(nn.Module):
    def __init__(self, channels, reduction=4):
        super().__init__()
        reduced = max(1, channels // reduction)
        self.fc1 = nn.Linear(channels, reduced)
        self.fc2 = nn.Linear(reduced, channels)

    def forward(self, x):
        b, c, _, _ = x.shape
        y = x.mean(dim=[2, 3]) # Global avg pool
        y = torch.relu(self.fc1(y))
        y = torch.sigmoid(self.fc2(y))
        return x * y.view(b, c, 1, 1)

print("\nSqueeze-Excitation:")
print(" - Global average pool to get channel statistics")
print(" - FC layers to learn channel importance")
print(" - Sigmoid to get attention weights")
print(" - Multiply to reweight channels")

print("\nActivation: SiLU (Swish)")
print(" SiLU(x) = x * sigmoid(x)")

MBConv (Mobile Inverted Bottleneck)
=====
MBConv structure:
1. Expand: 1x1 conv (increase channels)
2. Depthwise: 3x3 or 5x5 depthwise conv
3. Squeeze-Excitation: Channel attention
4. Project: 1x1 conv (reduce channels)

Squeeze-Excitation:
- Global average pool to get channel statistics
- FC layers to learn channel importance
- Sigmoid to get attention weights
- Multiply to reweight channels

Activation: SiLU (Swish)
SiLU(x) = x * sigmoid(x)

```

## Summary

### Key Takeaways:

- Skip connections solve the degradation problem
- Residual blocks learn  $F(x) = H(x) - x$
- Bottleneck design reduces parameters
- EfficientNet scales depth, width, resolution together
- MBConv uses depthwise separable convolutions
- Squeeze-Excitation provides channel attention

### Practice Exercises:

1. Implement BasicBlock and Bottleneck
2. Train ResNet-18 on CIFAR-10
3. Compare ResNet vs VGG training
4. Visualize gradient flow with skip connections
5. Load and analyze EfficientNet