

## **Text Processing, Embeddings, and Word2Vec**

### **Lecture Overview**

This lecture introduces natural language processing fundamentals including tokenization, word embeddings, and the Word2Vec algorithm.

Topics Covered:

- Text preprocessing pipeline
- One-hot vs dense embeddings
- Word2Vec: Skip-gram and CBOW
- Pre-trained embeddings
- Simple text classification

# 1. Text Preprocessing

## 1.1 Text to Numbers

```
import torch
import torch.nn as nn
import numpy as np

print("Text Preprocessing")
print("=*70)

print("Challenge: Neural networks need numbers, not text!")

print("\nPreprocessing Pipeline:")
print(" 1. Tokenization: Split text into tokens")
print(" 2. Vocabulary: Map tokens to integers")
print(" 3. Embedding: Map integers to dense vectors")

# Simple tokenization
text = "The cat sat on the mat."
tokens = text.lower().replace(".", "").split()
print(f"\nText: {text}")
print(f"Tokens: {tokens}")

# Build vocabulary
vocab = {"<PAD>": 0, "<UNK>": 1}
for token in tokens:
    if token not in vocab:
        vocab[token] = len(vocab)

print(f"\nVocabulary: {vocab}")

# Convert to indices
indices = [vocab.get(t, vocab["<UNK>"]) for t in tokens]
print(f"Indices: {indices}")

print("\nTokenization approaches:")
print(" - Word-level: 'playing' -> ['playing']")
print(" - Subword (BPE): 'playing' -> ['play', 'ing']")
print(" - Character: 'cat' -> ['c', 'a', 't']")

Text Preprocessing
=====
Challenge: Neural networks need numbers, not text!

Preprocessing Pipeline:
 1. Tokenization: Split text into tokens
 2. Vocabulary: Map tokens to integers
 3. Embedding: Map integers to dense vectors

Text: 'The cat sat on the mat.'
Tokens: ['the', 'cat', 'sat', 'on', 'the', 'mat']

Vocabulary: {'<PAD>': 0, '<UNK>': 1, 'the': 2, 'cat': 3, 'sat': 4, 'on': 5, 'mat': 6}
Indices: [2, 3, 4, 5, 2, 6]

Tokenization approaches:
 - Word-level: 'playing' -> ['playing']
 - Subword (BPE): 'playing' -> ['play', 'ing']
 - Character: 'cat' -> ['c', 'a', 't']
```

## 2. Word Embeddings

### 2.1 One-Hot vs Dense Embeddings

```
print("One-Hot vs Dense Embeddings")
print("=*70)

vocab_size = 10000

print("One-Hot Encoding:")
print(f"  Vector size: {vocab_size}")
print(f"  'cat' = [0, 0, ..., 1, ..., 0, 0]")
print("  Problems:")
print("    - Huge vectors (10,000+ dimensions)")
print("    - No similarity: cat · dog = 0")
print("    - Sparse and inefficient")

print("\nDense Embeddings:")
embed_dim = 300
print(f"  Vector size: {embed_dim}")
print(f"  'cat' = [0.2, -0.5, 0.8, ...]")
print("  Benefits:")
print("    - Compact representation")
print("    - Captures semantic similarity")
print("    - cat · dog > 0 (both animals)")

# PyTorch embedding layer
embedding = nn.Embedding(vocab_size, embed_dim)
print(f"\nnn.Embedding({vocab_size}, {embed_dim})")
print(f"  Weight shape: {embedding.weight.shape}")
print(f"  Parameters: {embedding.weight.numel():,}")

# Lookup
indices = torch.tensor([42, 100, 5])
vectors = embedding(indices)
print(f"\nLookup indices {indices.tolist()}")
print(f"  Output shape: {vectors.shape}")

One-Hot vs Dense Embeddings
=====
One-Hot Encoding:
  Vector size: 10000
  'cat' = [0, 0, ..., 1, ..., 0, 0]
  Problems:
    - Huge vectors (10,000+ dimensions)
    - No similarity: cat · dog = 0
    - Sparse and inefficient

Dense Embeddings:
  Vector size: 300
  'cat' = [0.2, -0.5, 0.8, ...]
  Benefits:
    - Compact representation
    - Captures semantic similarity
    - cat · dog > 0 (both animals)

nn.Embedding(10000, 300)
  Weight shape: torch.Size([10000, 300])
  Parameters: 3,000,000

Lookup indices [42, 100, 5]
  Output shape: torch.Size([3, 300])
```

### 2.2 Word2Vec

```
print("Word2Vec (Mikolov et al., 2013)")
print("=*70)

print("Key idea: Words with similar contexts have similar meanings")
print("  'The cat sat on the ___' -> mat, rug, floor")
print("  'The dog sat on the ___' -> mat, rug, floor")
print("  => cat and dog have similar embeddings")
```

```

print("\nTwo architectures:")
print("\n1. Skip-gram: Predict context from center word")
print("  Input: 'cat' -> Predict: 'the', 'sat', 'on'")
print("  Better for rare words")

print("\n2. CBOW (Continuous Bag of Words):")
print("  Input: 'the', 'sat', 'on' -> Predict: 'cat'")
print("  Faster training")

print("\nTraining objective (Skip-gram):")
print("  maximize P(context | center)")
print("  = softmax(v_context · v_center)")

print("\nFamous property - Word Analogies:")
print("  king - man + woman = queen")
print("  paris - france + italy = rome")

print("\nUsing pre-trained embeddings:")
print("  - Word2Vec (Google)")
print("  - GloVe (Stanford)")
print("  - FastText (Facebook)")

Word2Vec (Mikolov et al., 2013)
=====
Key idea: Words with similar contexts have similar meanings
'The cat sat on the ___' -> mat, rug, floor
'The dog sat on the ___' -> mat, rug, floor
=> cat and dog have similar embeddings

```

Two architectures:

1. Skip-gram: Predict context from center word  
Input: 'cat' -> Predict: 'the', 'sat', 'on'  
Better for rare words
2. CBOW (Continuous Bag of Words):  
Input: 'the', 'sat', 'on' -> Predict: 'cat'  
Faster training

Training objective (Skip-gram):  
maximize  $P(\text{context} \mid \text{center})$   
=  $\text{softmax}(v_{\text{context}} \cdot v_{\text{center}})$

Famous property - Word Analogies:  
king - man + woman = queen  
paris - france + italy = rome

Using pre-trained embeddings:  
- Word2Vec (Google)  
- GloVe (Stanford)  
- FastText (Facebook)

### 3. Text Classification Model

#### 3.1 Simple Text Classifier

```
print("Simple Text Classification")
print("*" * 70)

class TextClassifier(nn.Module):
    def __init__(self, vocab_size, embed_dim, num_classes):
        super().__init__()
        self.embedding = nn.Embedding(vocab_size, embed_dim, padding_idx=0)
        self.fc = nn.Linear(embed_dim, num_classes)

    def forward(self, x):
        # x: (batch, seq_len)
        embedded = self.embedding(x)  # (batch, seq_len, embed_dim)
        # Average pooling over sequence
        pooled = embedded.mean(dim=1)  # (batch, embed_dim)
        return self.fc(pooled)

model = TextClassifier(vocab_size=10000, embed_dim=128, num_classes=2)
print(model)

# Example: Sentiment classification
x = torch.randint(0, 10000, (4, 50))  # 4 texts, 50 tokens each
out = model(x)
print(f"\nInput shape: {x.shape}")
print(f"Output shape: {out.shape}")

print("\nPooling strategies:")
print(" - Mean: Average all embeddings")
print(" - Max: Max over each dimension")
print(" - [CLS]: Use special token (BERT-style)")

Simple Text Classification
=====
TextClassifier(
    (embedding): Embedding(10000, 128, padding_idx=0)
    (fc): Linear(in_features=128, out_features=2, bias=True)
)

Input shape: torch.Size([4, 50])
Output shape: torch.Size([4, 2])

Pooling strategies:
 - Mean: Average all embeddings
 - Max: Max over each dimension
 - [CLS]: Use special token (BERT-style)
```

## Summary

### Key Takeaways:

- Tokenization converts text to token sequences
- Dense embeddings capture semantic similarity
- Word2Vec learns embeddings from context prediction
- Similar words have similar embeddings
- Pre-trained embeddings provide good initialization
- Pooling converts variable-length to fixed-length

### Practice Exercises:

1. Implement word tokenization from scratch

2. Build vocabulary from corpus
3. Use pre-trained GloVe embeddings
4. Implement simple text classifier
5. Visualize word embeddings with t-SNE