

Hugging Face Transformers Library

Lecture Overview

This lecture covers the Hugging Face ecosystem including pipelines, tokenizers, models, and the Trainer API for fine-tuning.

Topics Covered:

- Hugging Face ecosystem
- Pipelines for quick inference
- Tokenizers and models
- AutoModel classes
- Trainer API for fine-tuning

1. Introduction to Hugging Face

1.1 The Hugging Face Ecosystem

```
print("Hugging Face Ecosystem")
print("="*70)

print("Hugging Face: The GitHub of ML")

print("\nKey components:")
print("  1. Transformers library: Pre-trained models")
print("  2. Datasets library: Standard datasets")
print("  3. Tokenizers library: Fast tokenization")
print("  4. Hub: Model and dataset repository")

print("\nWhy Hugging Face?")
print("  - 200,000+ pre-trained models")
print("  - Consistent API across models")
print("  - Easy to use pipelines")
print("  - Active community")

print("\nInstallation:")
print("  pip install transformers")
print("  pip install datasets")

print("\nSupported frameworks:")
print("  - PyTorch")
print("  - TensorFlow")
print("  - JAX/Flax")

Hugging Face Ecosystem
=====
Hugging Face: The GitHub of ML

Key components:
  1. Transformers library: Pre-trained models
  2. Datasets library: Standard datasets
  3. Tokenizers library: Fast tokenization
  4. Hub: Model and dataset repository

Why Hugging Face?
  - 200,000+ pre-trained models
  - Consistent API across models
  - Easy to use pipelines
  - Active community

Installation:
  pip install transformers
  pip install datasets

Supported frameworks:
  - PyTorch
  - TensorFlow
  - JAX/Flax
```

2. Pipelines

2.1 Using Pipelines

```
print("Hugging Face Pipelines")
print("="*70)

print("Pipelines: Easiest way to use models")
print("""
from transformers import pipeline

# Sentiment analysis
classifier = pipeline("sentiment-analysis")
result = classifier("I love this movie!")
# [{'label': 'POSITIVE', 'score': 0.9998}]

# Text generation
generator = pipeline("text-generation", model="gpt2")
result = generator("Once upon a time", max_length=50)

# Question answering
qa = pipeline("question-answering")
result = qa(question="What is AI?",
            context="AI is artificial intelligence...")

# Named entity recognition
ner = pipeline("ner")
result = ner("Microsoft was founded by Bill Gates")

# Summarization
summarizer = pipeline("summarization")
result = summarizer(long_text, max_length=100)
""")

print("Available pipelines:")
print(" - sentiment-analysis, text-classification")
print(" - text-generation, text2text-generation")
print(" - question-answering, fill-mask")
print(" - ner, translation, summarization")
print(" - zero-shot-classification")

Hugging Face Pipelines
=====
Pipelines: Easiest way to use models

from transformers import pipeline

# Sentiment analysis
classifier = pipeline("sentiment-analysis")
result = classifier("I love this movie!")
# [{'label': 'POSITIVE', 'score': 0.9998}]

# Text generation
generator = pipeline("text-generation", model="gpt2")
result = generator("Once upon a time", max_length=50)

# Question answering
qa = pipeline("question-answering")
result = qa(question="What is AI?",
            context="AI is artificial intelligence...")

# Named entity recognition
ner = pipeline("ner")
result = ner("Microsoft was founded by Bill Gates")

# Summarization
summarizer = pipeline("summarization")
result = summarizer(long_text, max_length=100)

Available pipelines:
- sentiment-analysis, text-classification
- text-generation, text2text-generation
- question-answering, fill-mask
- ner, translation, summarization
```

- zero-shot-classification

3. Tokenizers and Models

3.1 Tokenizers

```
print("Tokenizers")
print("*" * 70)

print("Loading tokenizers:")
print("""
from transformers import AutoTokenizer

tokenizer = AutoTokenizer.from_pretrained("bert-base-uncased")

# Tokenize text
text = "Hello, how are you?"
tokens = tokenizer(text)
print(tokens)
# {'input_ids': [101, 7592, 1010, 2129, 2024, 2017, 1029, 102],
#  'token_type_ids': [0, 0, 0, 0, 0, 0, 0, 0],
#  'attention_mask': [1, 1, 1, 1, 1, 1, 1, 1]}

# Decode back
decoded = tokenizer.decode(tokens['input_ids'])
print(decoded)  # '[CLS] hello, how are you? [SEP]'

# Batch tokenization with padding
texts = ["Hello", "How are you doing today?"]
batch = tokenizer(texts, padding=True, return_tensors="pt")
""")

print("\nKey tokenizer methods:")
print("  tokenizer(text) - Tokenize")
print("  tokenizer.decode(ids) - Decode")
print("  tokenizer.encode(text) - Get input_ids only")
print("  tokenizer.vocab_size - Vocabulary size")

Tokenizers
=====
Loading tokenizers:

from transformers import AutoTokenizer

tokenizer = AutoTokenizer.from_pretrained("bert-base-uncased")

# Tokenize text
text = "Hello, how are you?"
tokens = tokenizer(text)
print(tokens)
# {'input_ids': [101, 7592, 1010, 2129, 2024, 2017, 1029, 102],
#  'token_type_ids': [0, 0, 0, 0, 0, 0, 0, 0],
#  'attention_mask': [1, 1, 1, 1, 1, 1, 1, 1]}

# Decode back
decoded = tokenizer.decode(tokens['input_ids'])
print(decoded)  # '[CLS] hello, how are you? [SEP]'

# Batch tokenization with padding
texts = ["Hello", "How are you doing today?"]
batch = tokenizer(texts, padding=True, return_tensors="pt")

Key tokenizer methods:
  tokenizer(text) - Tokenize
  tokenizer.decode(ids) - Decode
  tokenizer.encode(text) - Get input_ids only
  tokenizer.vocab_size - Vocabulary size
```

3.2 Loading Models

```
print("Loading Models")
print("*" * 70)

print("AutoModel classes:")
print("")
```

```

from transformers import (
    AutoModel,
    AutoModelForSequenceClassification,
    AutoModelForQuestionAnswering,
    AutoModelForCausalLM,
)

# Base model (no head)
model = AutoModel.from_pretrained("bert-base-uncased")

# Classification
model = AutoModelForSequenceClassification.from_pretrained(
    "bert-base-uncased",
    num_labels=2
)

# Text generation
model = AutoModelForCausalLM.from_pretrained("gpt2")

# Get model outputs
outputs = model(**tokenizer(text, return_tensors="pt"))
# outputs.last_hidden_state, outputs.logits, etc.
"""

print("\nPopular models:")
print("  BERT: bert-base-uncased, bert-large-uncased")
print("  RoBERTa: roberta-base, roberta-large")
print("  DistilBERT: distilbert-base-uncased")
print("  GPT-2: gpt2, gpt2-medium, gpt2-large")
print("  T5: t5-small, t5-base, t5-large")

Loading Models
=====
AutoModel classes:

from transformers import (
    AutoModel,
    AutoModelForSequenceClassification,
    AutoModelForQuestionAnswering,
    AutoModelForCausalLM,
)

# Base model (no head)
model = AutoModel.from_pretrained("bert-base-uncased")

# Classification
model = AutoModelForSequenceClassification.from_pretrained(
    "bert-base-uncased",
    num_labels=2
)

# Text generation
model = AutoModelForCausalLM.from_pretrained("gpt2")

# Get model outputs
outputs = model(**tokenizer(text, return_tensors="pt"))
# outputs.last_hidden_state, outputs.logits, etc.

Popular models:
  BERT: bert-base-uncased, bert-large-uncased
  RoBERTa: roberta-base, roberta-large
  DistilBERT: distilbert-base-uncased
  GPT-2: gpt2, gpt2-medium, gpt2-large
  T5: t5-small, t5-base, t5-large

```

4. Training with Trainer

4.1 Trainer API

```
print("Hugging Face Trainer")
print("="*70)

print("Complete training workflow:")
print("""
from transformers import (
    AutoModelForSequenceClassification,
    AutoTokenizer,
    Trainer,
    TrainingArguments,
)
from datasets import load_dataset

# Load dataset
dataset = load_dataset("imdb")

# Load model and tokenizer
model = AutoModelForSequenceClassification.from_pretrained(
    "bert-base-uncased", num_labels=2
)
tokenizer = AutoTokenizer.from_pretrained("bert-base-uncased")

# Tokenize dataset
def tokenize(batch):
    return tokenizer(batch["text"], padding=True, truncation=True)

dataset = dataset.map(tokenize, batched=True)

# Training arguments
args = TrainingArguments(
    output_dir='./results',
    num_train_epochs=3,
    per_device_train_batch_size=16,
    learning_rate=2e-5,
    evaluation_strategy="epoch",
)

# Create trainer
trainer = Trainer(
    model=model,
    args=args,
    train_dataset=dataset["train"],
    eval_dataset=dataset["test"],
)

# Train!
trainer.train()
""")

Hugging Face Trainer
=====
Complete training workflow:

from transformers import (
    AutoModelForSequenceClassification,
    AutoTokenizer,
    Trainer,
    TrainingArguments,
)
from datasets import load_dataset

# Load dataset
dataset = load_dataset("imdb")

# Load model and tokenizer
model = AutoModelForSequenceClassification.from_pretrained(
    "bert-base-uncased", num_labels=2
)
tokenizer = AutoTokenizer.from_pretrained("bert-base-uncased")
```

```

# Tokenize dataset
def tokenize(batch):
    return tokenizer(batch["text"], padding=True, truncation=True)

dataset = dataset.map(tokenize, batched=True)

# Training arguments
args = TrainingArguments(
    output_dir='./results',
    num_train_epochs=3,
    per_device_train_batch_size=16,
    learning_rate=2e-5,
    evaluation_strategy="epoch",
)

# Create trainer
trainer = Trainer(
    model=model,
    args=args,
    train_dataset=dataset["train"],
    eval_dataset=dataset["test"],
)

# Train!
trainer.train()

```

Summary

Key Takeaways:

- Hugging Face provides unified API for transformers
- Pipelines are easiest way to use models
- AutoTokenizer/AutoModel load correct classes
- Always use padding and truncation
- Trainer simplifies training workflow
- Hub has 200,000+ pre-trained models

Practice Exercises:

1. Use pipeline for sentiment analysis
2. Tokenize text with BERT tokenizer
3. Fine-tune BERT on custom dataset
4. Compare different pre-trained models
5. Push trained model to Hub